

Serverless Architecture

Mischa Spiegelmock

Co-Founder,  ET BRIDGE

About Me

- From Berkeley, California
- Co-founder of JetBridge in SF/Kyiv/Wrocław
- Do computer shit
- Been doing computer shit a long time
- I'm a **developer**
- AWS Certified Solutions Architect

This Talk

- Ops from a developer's point of view
- Cloud-native application architecture
- Pros/cons of serverless applications
- Costs of Serverless

Ops

As a developer



What's the ?

Just somebody else's
computers?

- Use economies of scale of cloud providers
- Worry about higher-level issues more related to business
- Increase developer “leverage” 💪
- Don’t solve already solved problems (e.g. deploying web applications)







Serverless

Mindset

Business Value

Functions As A Service

“FaaS”

AWS Lambda



FaaS

- Write a function
- Deploy it “somewhere”
- Invoke it “somehow”

Write A Function

How it works

Run



```
1 exports.handler = (event, context, callback) => {  
2     // Succeed with the string "Hello world!"  
3     callback(null, 'Hello 🌎!');  
4 };
```

Deploy It Somewhere

```
Last login: Mon Feb 25 17:12:46 on ttys015
```

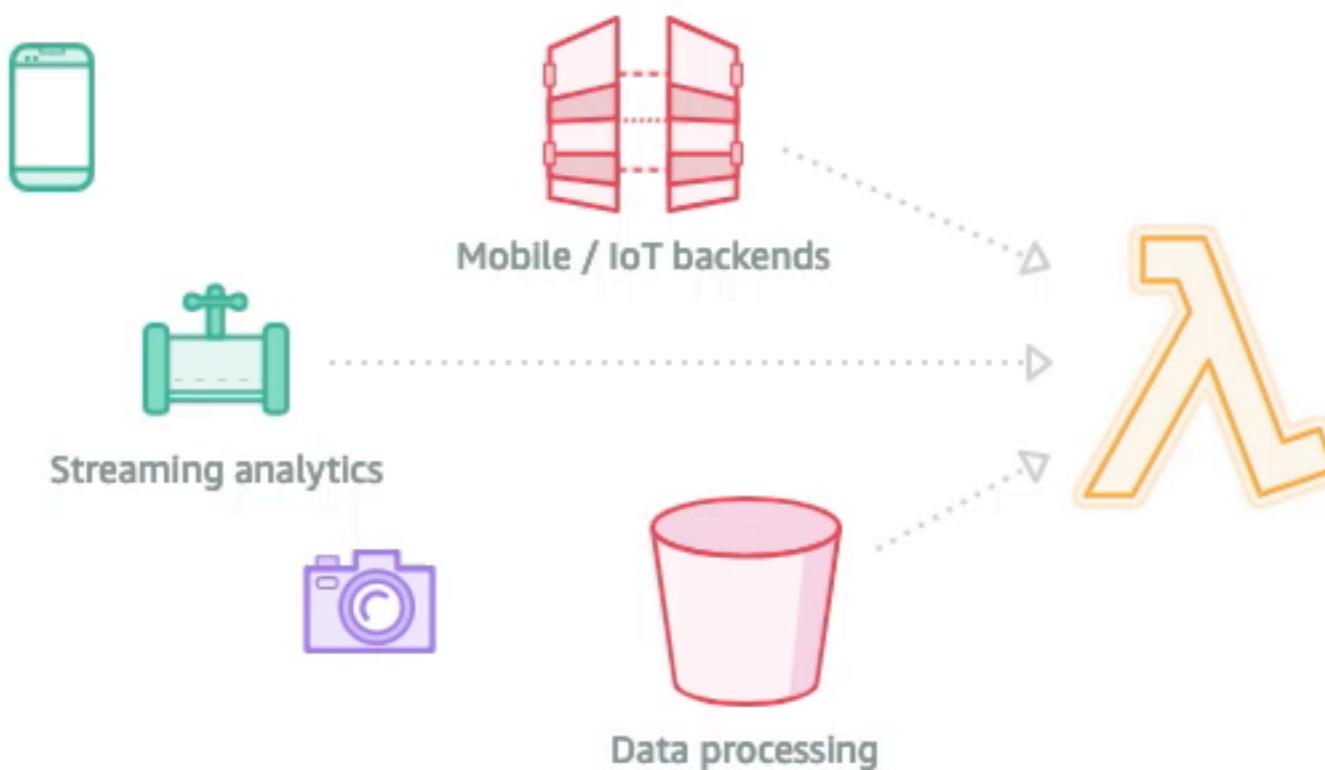
```
→ ~ slsp
```

```
Last login: Thu Feb 21 18:04:44 on console
```

```
→ crmservice git:(heroku-single-project) ✘
```

Invoke It Somehow

How it works

[Previous](#)[Next: Scale seamlessly](#)

Lambda responds to events

Once you create Lambda functions, you can configure them to respond to events from a variety of sources. Try sending a mobile notification, streaming data to Lambda, or placing a photo in an S3 bucket.

So what?

For Free:

High Availability

Not My Problem

Versioning

Every function gets a version

Autoscaling

1,000 polling concurrent invocations

Free-ish

Free as in 

Function-Level Permissions

Least privilege resource access

Filter events

all 2019-02-24 (19:38:23)

Time (UTC +00:00)	Message
2019-02-25	
▶ 19:24:34	START RequestId: 67c09671-9369-432e-9762-74e47662cfb9 Version: \$LATEST
▶ 19:24:34	END RequestId: 67c09671-9369-432e-9762-74e47662cfb9
▶ 19:24:34	REPORT RequestId: 67c09671-9369-432e-9762-74e47662cfb9 Duration: 38.15 ms Billed Duration: 100 ms Memory Size: 128 MB Max Memory Used: 41 MB
▶ 19:24:34	START RequestId: bce0fba0-9390-4218-9a86-9557ba8f046c Version: \$LATEST
▶ 19:24:34	END RequestId: bce0fba0-9390-4218-9a86-9557ba8f046c
▶ 19:24:34	REPORT RequestId: bce0fba0-9390-4218-9a86-9557ba8f046c Duration: 12.40 ms Billed Duration: 100 ms Memory Size: 128 MB Max Memory Used: 41 MB
▶ 19:24:34	START RequestId: 942d2d7b-71fe-4f86-be43-bc449b0d0bb3 Version: \$LATEST
▶ 19:24:34	END RequestId: 942d2d7b-71fe-4f86-be43-bc449b0d0bb3
▶ 19:24:34	REPORT RequestId: 942d2d7b-71fe-4f86-be43-bc449b0d0bb3 Duration: 20.35 ms Billed Duration: 100 ms Memory Size: 128 MB Max Memory Used: 41 MB
▶ 19:24:42	START RequestId: d9d89820-fcd2-4178-bbc3-65dccde83dd5 Version: \$LATEST
▶ 19:24:42	END RequestId: d9d89820-fcd2-4178-bbc3-65dccde83dd5

ALARM 0 INSUFFICIENT DATA 0 OK 0

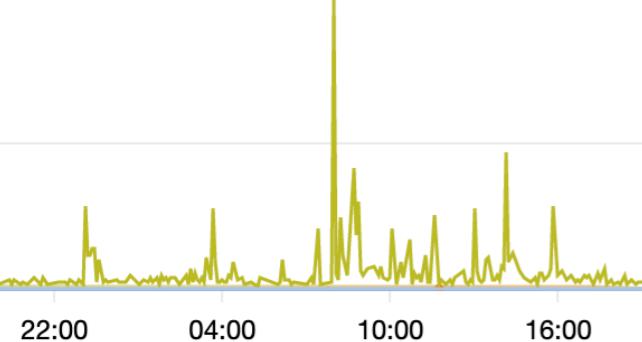
Invocations Sum

Various units

1.27k

635

1.00



- jb-contact-form-dev-unitFact...
- MakePyPIDep
- fingerprint-custom-domain-dev...
- salesPipe-prod-dispatch
- salesPipe-dev-handleSlackRe...
- costimator-prod-submit
- check_worker_queue_size
- monitor

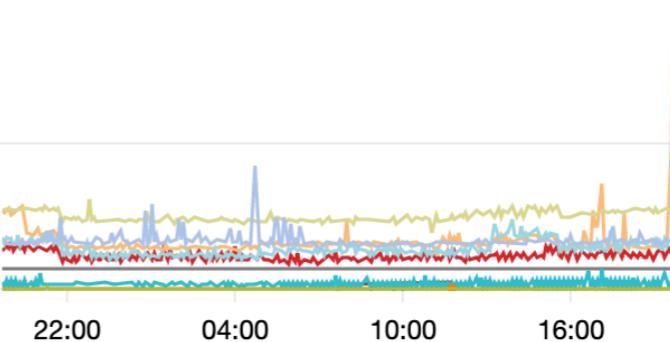
Duration Average

Various units

36.0k

18.0k

0.862



- jb-contact-form-dev-unitFact...
- MakePyPIDep
- fingerprint-custom-domain-dev...
- salesPipe-prod-dispatch
- salesPipe-dev-handleSlackRe...
- costimator-prod-submit
- check_worker_queue_size
- monitor

Errors Sum

Various units

3.00

1.50



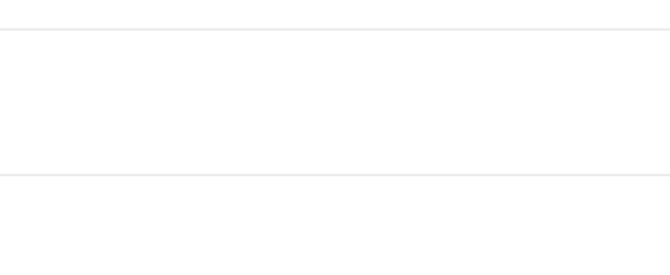
- jb-contact-form-dev-unitFact...
- MakePyPIDep
- fingerprint-custom-domain-dev...
- salesPipe-prod-dispatch
- salesPipe-dev-handleSlackRe...
- costimator-prod-submit
- check_worker_queue_size
- monitor

Throttles Sum

Various units

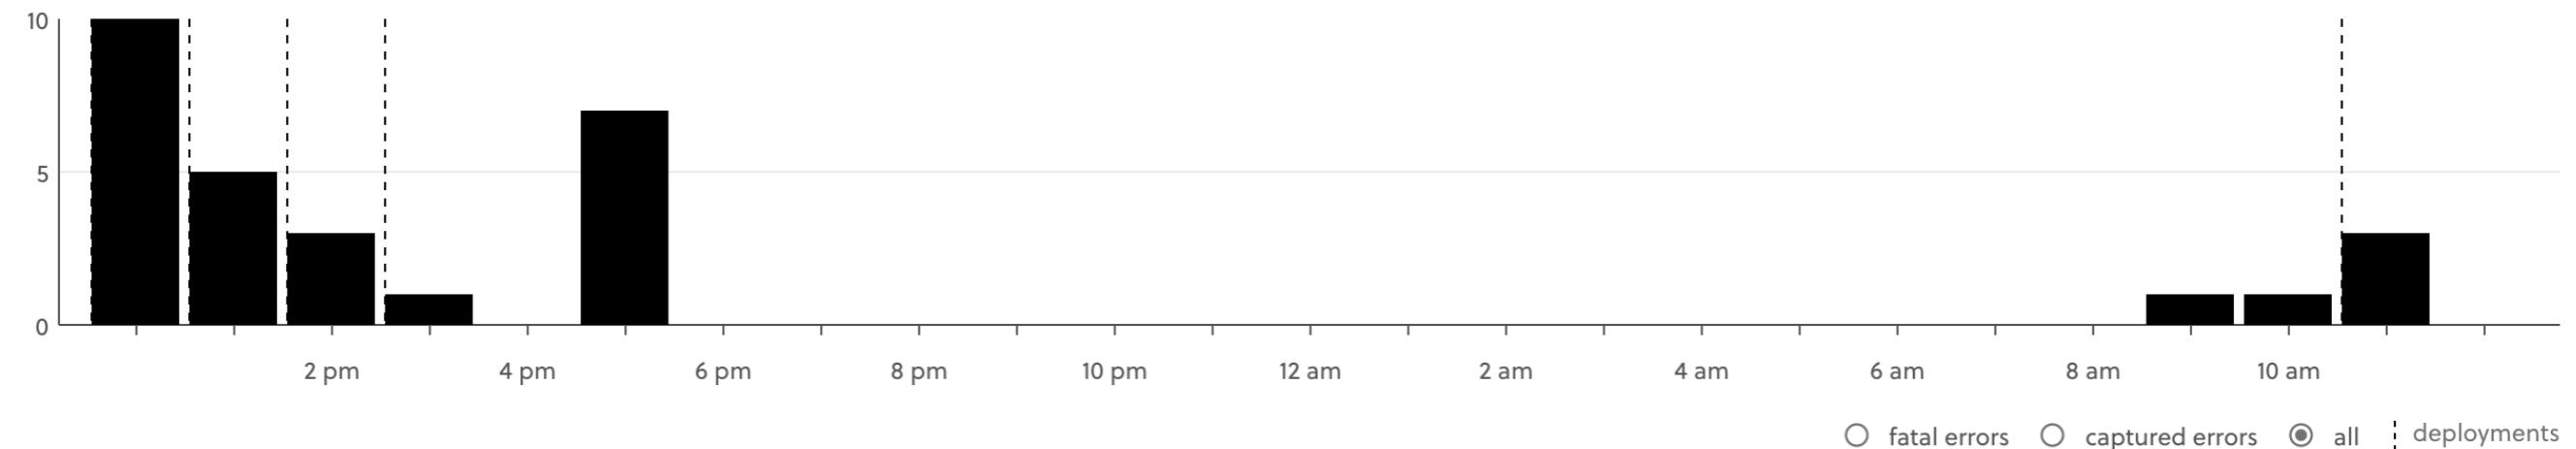
1.00

0.5



- jb-contact-form-dev-unitFact...
- MakePyPIDep
- fingerprint-custom-domain-dev...
- salesPipe-prod-dispatch
- salesPipe-dev-handleSlackRe...
- costimator-prod-submit
- check_worker_queue_size
- monitor

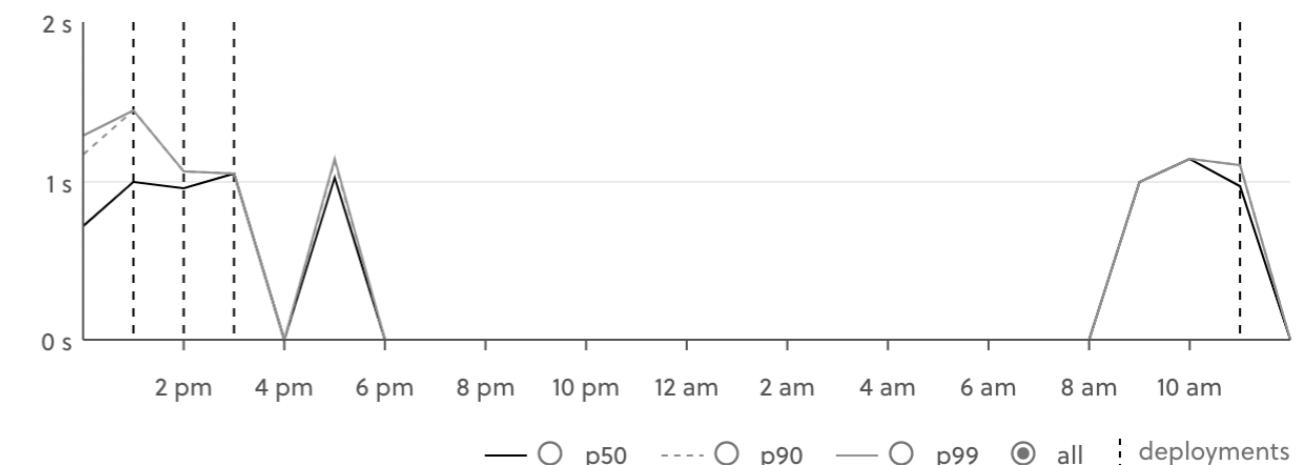
invocations & errors



memory usage



durations



request

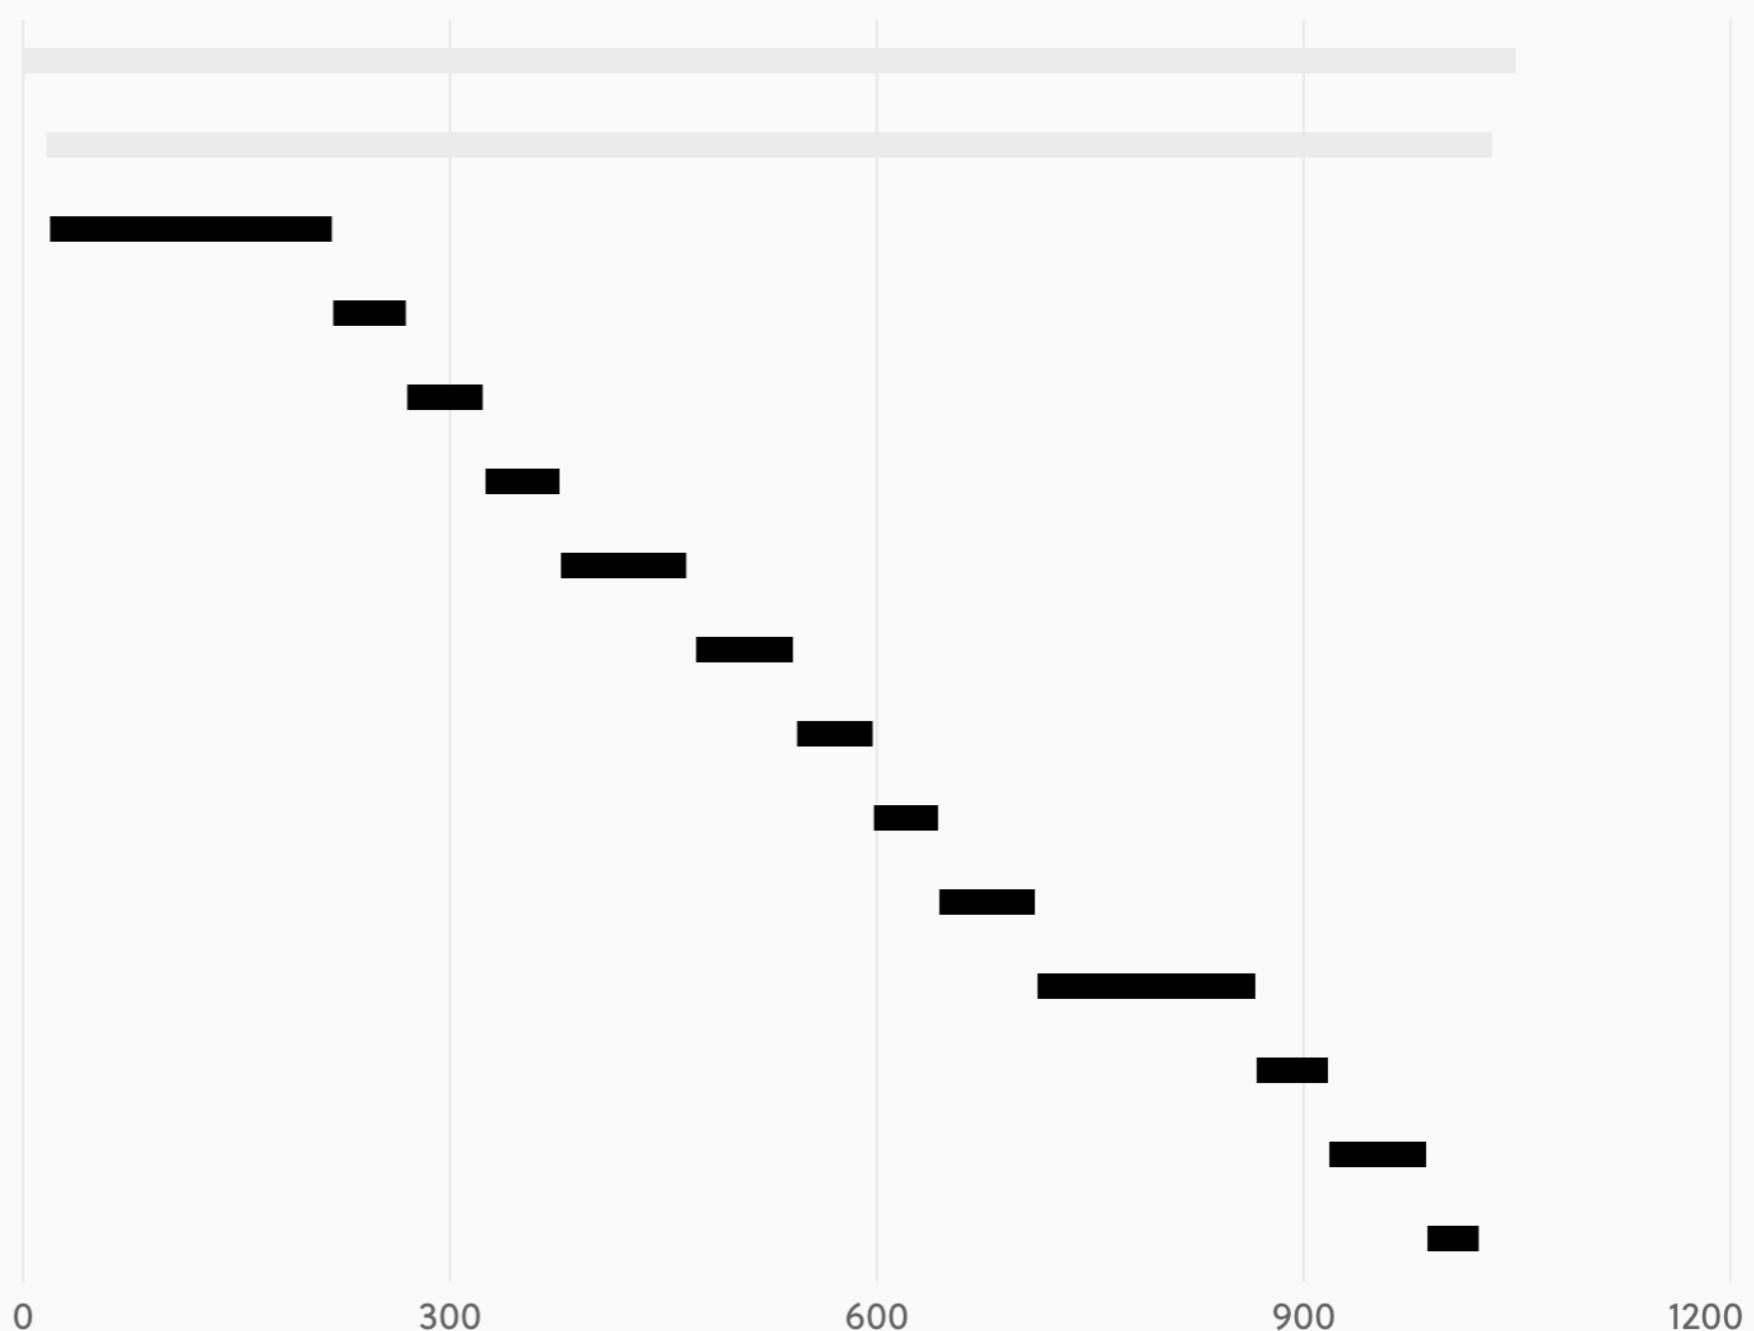
endpoint /{proxy+}
method POST
status 200
message -
latency 1 s

spans

DURATION OPERATION

1 s	total duration
1 s	function invocation
198 ms	CreateKeysAndCertificate iot
51 ms	CreateThing iot
53 ms	AttachPolicy iot
52 ms	AttachThingPrincipal iot
88 ms	CreateCoreDefinition greengrass
68 ms	DescribeStacks cloudformation
53 ms	GetFunctionDefinition greengrass
45 ms	GetSubscriptionDefinition greengrass
67 ms	GetLoggerDefinition greengrass
153 ms	CreateGroup greengrass
50 ms	AssociateRoleToGroup greengrass
68 ms	UpdateThingShadow iot-data
36 ms	DescribeEndpoint iot

SPANS



/ days ▾

view api requests ▾

// ×

/{proxy+} ×

all status codes



all methods



duration ms sec min 0



to 15

▼

requests: 89

DATE & TIME	METHOD	ENDPOINT	DURATION	STATUS/MESSAGE
feb 12 11:35 am	POST	/{proxy+}	928	● 200: ok
feb 12 11:35 am	POST	/{proxy+}	2486	● 200: ok
feb 12 11:26 am	POST	/{proxy+}	2395	● 500: internal server error
feb 12 10:10 am	POST	/{proxy+}	2539	● 200: ok
feb 12 9:34 am	POST	/{proxy+}	2370	● 200: ok
feb 11 5:43 pm	POST	/{proxy+}	885	● 200: ok
feb 11 5:38 pm	POST	/{proxy+}	896	● 200: ok
feb 11 5:38 pm	POST	/{proxy+}	891	● 200: ok
feb 11 5:35 pm	POST	/{proxy+}	1040	● 200: ok
feb 11 5:34 pm	POST	/{proxy+}	2511	● 200: ok

↳ general statistics

total invocations

8k

cost total

\$ 0.05

total errors

377

billed duration

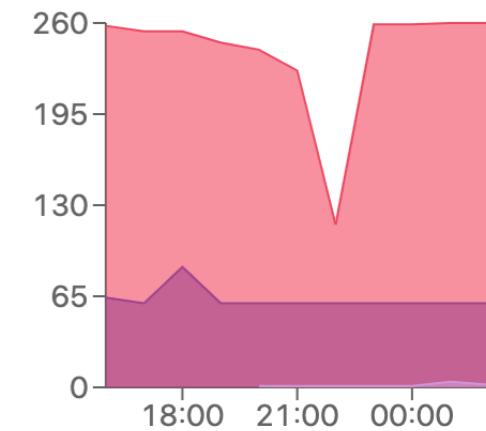
394 min

↳ cost

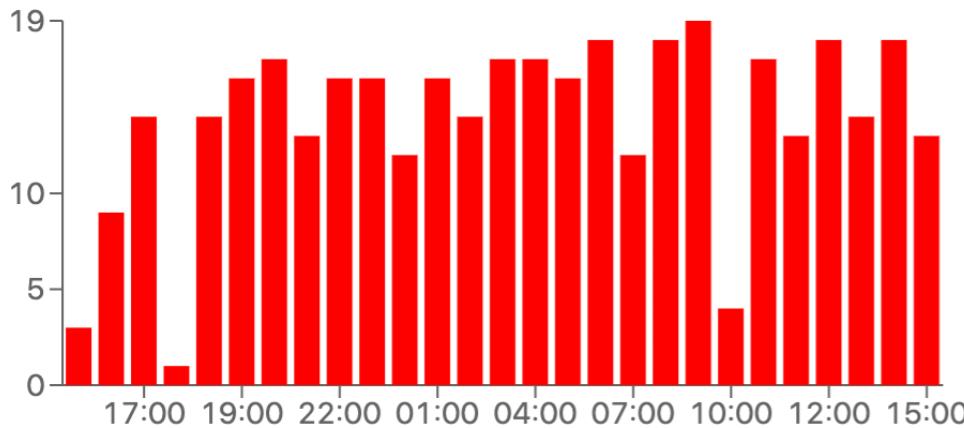
\$4
\$3
\$2
\$1
\$0

18:00 21:00 00:00 03:00 06:00 09:00 12:00 15:00

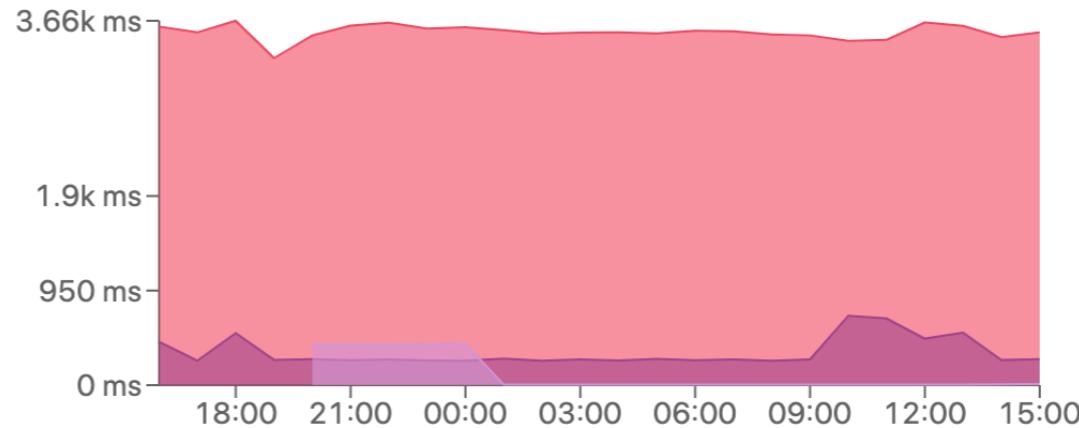
↳ total invocations



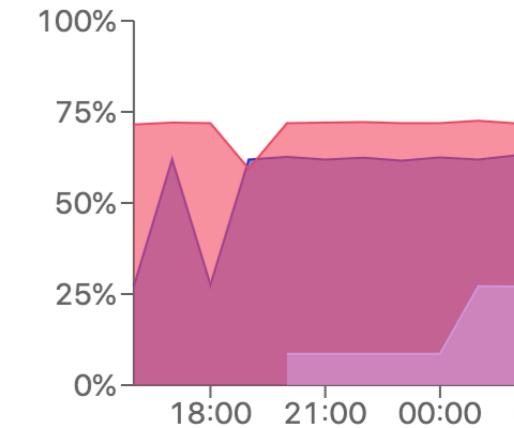
↳ invocation errors



↳ average durations



↳ memory utilization



Building Blocks

- AWS Lambda
- Google Cloud Functions
- Kubeless
- Azure
- IBM



serverless framework

The easy and open way to build serverless applications



Serverless (AWS)

- Interface to CloudWatch Logs (logging)
- Interface to CloudFormation (IaC, deployments)
- Interface to API Gateway (HTTP)
- Interface to Parameter Store/Secrets Manager (secrets)
- Plugins

Additional Building Blocks

- API Gateway (HTTP)
- CloudFront (CDN, TLS)
- RDS, Aurora Serverless, DynamoDB, AppSync (SQL, NoSQL, GraphQL)
- Cognito (Authentication)
- IoT (Greengrass)
- SNS/SES (SMS/Email)
- Secrets Manager
- ...

Benefits To FaaS

- Deployments: solved
- Infrastructure packaged with application
- Autoscaling
- High availability
- Tight integration with other building blocks
- No DevOps, just Dev

Downsides To FaaS

- Vendor lock-in
- Raw power
- Shared hosting
- Local development environment
- Cold start time

FaaS Use Cases

- Microservices
- Triggers (IoT, Storage, Message Queue, Auth, Alarms..)
- Web Applications
- Webhook Processors
- Streams (WebSockets, DynamoDB, Kinesis)

FaaS Is Not For

- Raw compute power (CPU, RAM, GPU, Disk)
- Long-running tasks
- Streaming media
- Large (disk size) applications
- High startup time
- Non-linux/amd64
- Dedicated hosting requirements (compliance)
- Real-time
- Probably a lot of other scenarios

Something In The
Middle?

AWS Fargate

(managed containers)

Costs

GB/s

+

Invocations

	Price
Requests	\$0.20 per 1M requests
Duration	\$0.000016667 for every GB-second
Memory (MB)	Price per 100ms
128	\$0.000000208
512	\$0.000000833
1024	\$0.000001667

E.g. 30,000,000 Invocations at 128MB RAM for 200ms

$$\begin{aligned}
 &= \\
 \$11.63
 \end{aligned}$$

Costs of Serverless

- Developer time
- Monitoring services
- API Gateway HTTP calls
- Lambda invocations
- Database

Savings of Serverless

- DevOps
- SiteOps
- Capital expenses (CAPEX)
- Unused server or database capacity
- Monitoring and maintenance

Future Of DevOps



The End

Email me: mischa@jetbridge.com

mish.dev

spiegelmock.com

jetbridge.com